

Lec04-2. Model-based Collaborative Filtering


contents

1. Basic Latent Factor Models
 2. Incorporating User and Item Biases
 3. Incorporating Implicit Feedback
-

1. Basic Latent Factor Models

Dimensionality Reduction

- Given data in a **high D -dimensional space**, we project data into a **low d -dimensional subspace**.



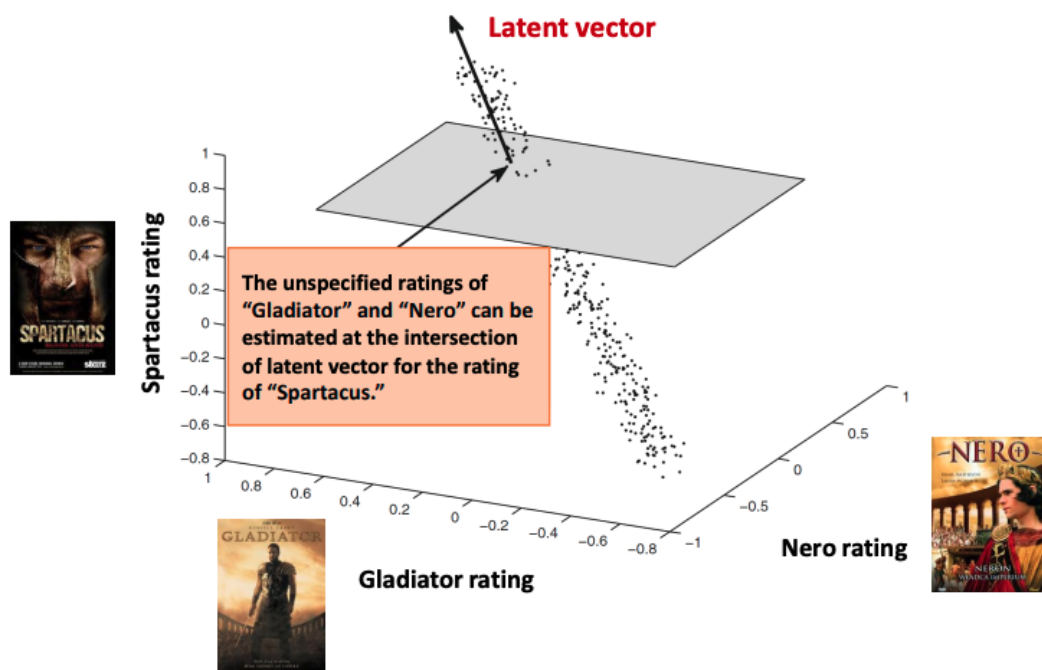
- How to represent **the axes of subspace** effectively?
 - Essential to preserve important features of data.

Why is Dimensionality Reduction?

- Key observation : **substantial portions of rows and columns** in the rating matrix are **highly correlated**.
 - The rating matrix has **redundant rows and columns**.
 - The **fully specified low-rank approximation** can be determined with a **small subset of the entries** in the original matrix.
- Dimensionality reduction is used to **rotate the axis to remove redundant pairwise correlations**.
 - The axis represents a **latent factor** for users and items.

Geometric View : Latent Factor Models

- The **ratings of the three movies** are highly **correlated**.
 - It can be arranged along a **1-dimensional line**.



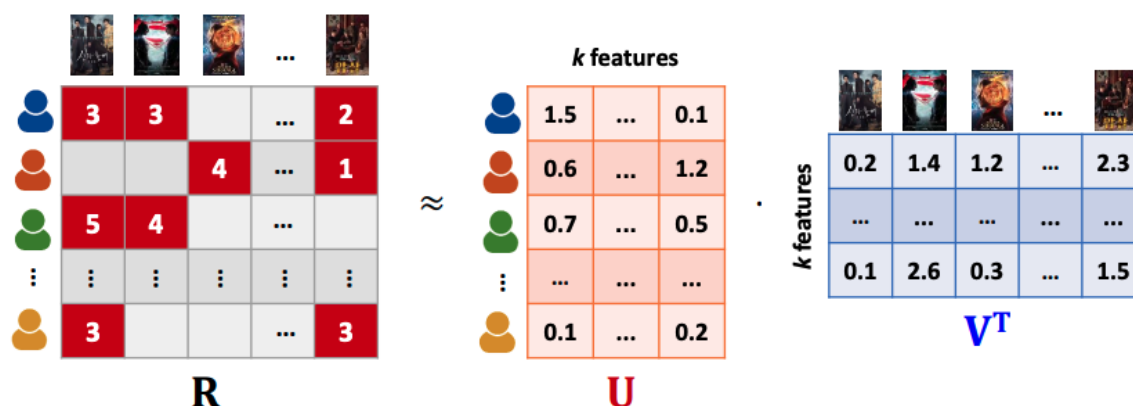
What are Latent Factor Models?

- Find a set of latent vectors, **minimizing the distance of user ratings from the hyperplane defined by the latent factors**.
 - **Capture the underlying redundancies** in the correlation structure of the data and **reconstruct all the missing values in one shot**.

- Note : If the data do not have any correlations or redundancies, the latent factor model does not work well.
- Matrix factorization is **a general solution to approximate a given matrix** for dimensionality reduction.
 - The latent vectors are **not always mutually orthogonal**.
 - SVD is the representative method of matrix factorization, in which basis vectors are orthogonal to each other.

What is Matrix Factorization?

- Given a matrix $R \in R^{m \times n}$, it can be approximately expressed in the product of low rank- k factors, $k \ll \min(m, n)$.
 - **U : latent user matrix (m x k matrix)**
 - Each user is represented by a latent vector(1 x k vector).
 - **V : latent item matrix (n x k matrix)**
 - Each item is represented by a latent vector(1 x k vector).



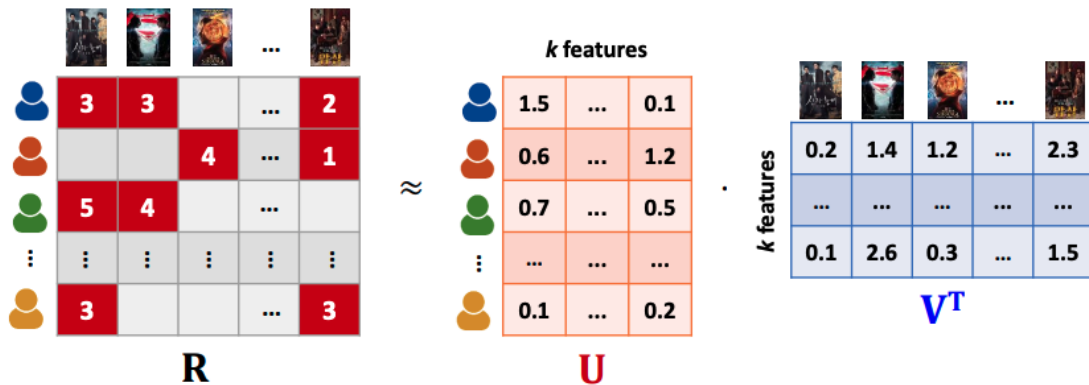
Objective Function for MF

- Factorize a matrix R into two latent matrices U and V .
 - The matrix R is approximated as a product of UV^T
 - **Note : missing ratings are ignored for model training.**

$$\min_{\mathbf{U}, \mathbf{V}} \sum_{u=1}^m \sum_{i=1}^n y_{ui} (r_{ui} - \mathbf{u}_u \mathbf{v}_i^T)^2$$

$$y_{ui} = \begin{cases} 1 & \text{if } r_{ui} \text{ exists} \\ 0 & \text{otherwise} \end{cases}$$


\mathbf{u}_u : u -th row of \mathbf{U} \mathbf{v}_i : i -th row of \mathbf{V}



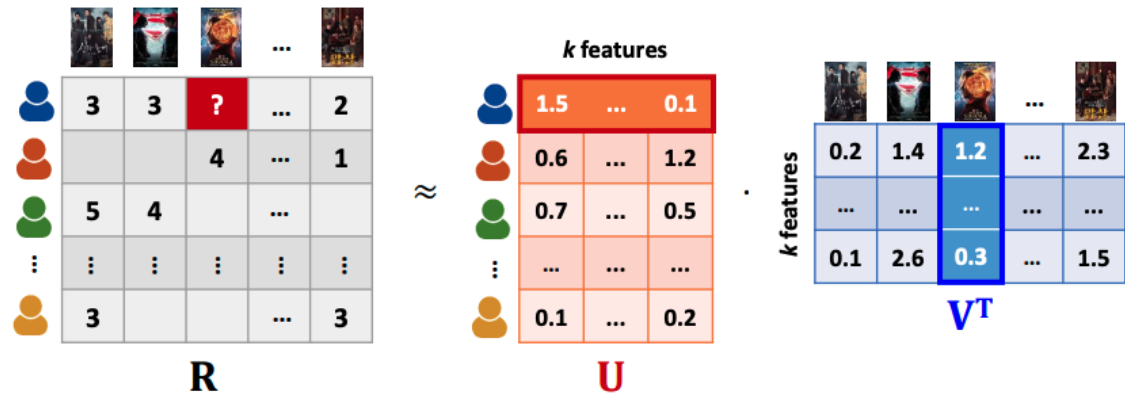
Conceptual View: Latent User Vectors

➤ Representing users in a **latent feature space**

- ◆ Note: two factors may not be orthogonal.



$$\hat{r}_{ui} = \mathbf{u}_u \mathbf{v}_i^T$$



Unconstrained Matrix Factorization

- Formulate an optimization problem for \mathbf{U} and \mathbf{V} .
 - There are no constraints on two matrices \mathbf{U} and \mathbf{V} .

$$\min \frac{1}{2} \sum_{(u,i) \in \mathcal{S}} e_{ui}^2 = \frac{1}{2} \sum_{(u,i) \in \mathcal{S}} (r_{ui} - \mathbf{u}_u \mathbf{v}_i^T)^2$$

\mathcal{S} : a set of observed user-item pairs in \mathbf{R}

- Minimize **the sum of the squared error** between the observed value and the predicted value for the entry (u, i) .
 - $e_{ui} = r_{ui} - \hat{r}_{ui}$, where $\hat{r}_{ui} = \mathbf{u}_u \mathbf{v}_i^T$

Recap: Gradient Descent (GD)

learning rate:
controlling the step size

Randomly initialize parameters \mathbf{w}^0 ,

Repeat

$$\mathbf{w}^{t+1} = \mathbf{w}^t - \alpha \left. \frac{dE}{d\mathbf{w}} \right|_{\mathbf{w}=\mathbf{w}^t}$$

Until the **stopping condition** is satisfied

- Fixed number of iterations
- $|E(\mathbf{w}^{t+1}) - E(\mathbf{w}^t)|$ is very small.

Computing Partial Derivatives



➤ How to compute the partial derivative of E for u_{uq} and v_{iq} ?

$$\begin{aligned}\frac{\partial E}{\partial u_{uq}} &= \sum_{i:(u,i) \in \mathcal{S}} \left(r_{ui} - \sum_{s=1}^k u_{us} \cdot v_{is} \right) (-v_{iq}) \quad \forall u \in \{1, \dots, m\}, q \in \{1, \dots, k\} \\ &= \sum_{i:(u,i) \in \mathcal{S}} (e_{ui}) (-v_{iq}) \quad \forall u \in \{1, \dots, m\}, q \in \{1, \dots, k\}\end{aligned}$$

$$\begin{aligned}\frac{\partial E}{\partial v_{iq}} &= \sum_{u:(u,i) \in \mathcal{S}} \left(r_{ui} - \sum_{s=1}^k u_{us} \cdot v_{is} \right) (-u_{uq}) \quad \forall i \in \{1, \dots, n\}, q \in \{1, \dots, k\} \\ &= \sum_{u:(u,i) \in \mathcal{S}} (e_{ui}) (-u_{uq}) \quad \forall i \in \{1, \dots, n\}, q \in \{1, \dots, k\}\end{aligned}$$

Training with GD

➤ **Input:** rating matrix **R** and learning rate α

Randomly initialize two matrices **U and **V**.**

Let $\mathcal{S} = \{(u, i): r_{ui} \text{ is observed}\}$.

Repeat

Compute each error e_{ui} as the observed entry of $\mathbf{R} - \mathbf{UV}^T$.

For each user-component pair (u, q) do

$$u_{uq}^{(t+1)} = u_{uq}^{(t)} - \alpha \sum_{i:(u,i) \in \mathcal{S}} e_{ui} \cdot (-v_{iq})$$

For each item-component pair (i, q) do

$$v_{iq}^{(t+1)} = v_{iq}^{(t)} - \alpha \sum_{u:(u,i) \in \mathcal{S}} e_{ui} \cdot (-u_{uq})$$

Until the **stopping condition is satisfied**

Training with GD



- How to perform the updates using matrix representations?
- Compute the error matrix $\mathbf{E} = \mathbf{R} - \mathbf{UV}^T$.
 - ◆ The unobserved entries of \mathbf{E} are set to 0.
- The updates can be computed as follows.

$$\begin{aligned}\mathbf{U} &= \mathbf{U} + \alpha \mathbf{E} \mathbf{V} \\ \mathbf{V} &= \mathbf{V} + \alpha \mathbf{E}^T \mathbf{U}\end{aligned}$$



- ◆ It can be executed to convergence.

Training with SGD



- It is specific to the observed entry $(u, i) \in \mathcal{S}$.
 - ◆ Update the relevant of $2k$ entries rather than $(mk + nk)$ entries.

$$\begin{aligned}e_{ui} &= r_{ui} - \mathbf{u}_u \mathbf{v}_i^T \\ \mathbf{u}_{uq} &\leftarrow \mathbf{u}_{uq} + \alpha e_{ui} \mathbf{v}_{iq} \quad \forall q \in \{1, \dots, k\} \\ \mathbf{v}_{iq} &\leftarrow \mathbf{v}_{iq} + \alpha e_{ui} \mathbf{u}_{uq} \quad \forall q \in \{1, \dots, k\}\end{aligned}$$

- It can be rewritten in vectorized form.

$$\begin{aligned}\mathbf{u}_u &\leftarrow \mathbf{u}_u + \alpha e_{ui} \mathbf{v}_i \\ \mathbf{v}_i &\leftarrow \mathbf{v}_i + \alpha e_{ui} \mathbf{u}_u\end{aligned}$$

- SGD is preferable when the data size is very large.

Adding Regularization Terms



➤ For model training, we introduce two terms.

- ◆ The goodness of fit is to reduce **the prediction error**.
- ◆ The regularization term is used to alleviate the **overfitting** problem.

Least square problems

$$\min \frac{1}{2} \sum_{(u,i) \in \mathcal{S}} e_{ui}^2 = \frac{1}{2} \sum_{(u,i) \in \mathcal{S}} (r_{ui} - \mathbf{u}_u \mathbf{v}_i^T)^2$$



Goodness of fit

Regularization

$$\min \frac{1}{2} \sum_{(u,i) \in \mathcal{S}} (r_{ui} - \mathbf{u}_u \mathbf{v}_i^T)^2 + \frac{\lambda}{2} \left(\sum_{u=1}^m \|\mathbf{u}_u\|^2 + \sum_{i=1}^n \|\mathbf{v}_i\|^2 \right)$$

Computing Partial Derivatives



➤ Interestingly, we can obtain almost the same results except for two terms $\lambda \mathbf{u}_{uq}$ and $\lambda \mathbf{v}_{iq}$.

$$\frac{\partial E}{\partial \mathbf{u}_{uq}} = \sum_{i:(u,i) \in \mathcal{S}} \left(r_{ui} - \sum_{s=1}^k \mathbf{u}_{us} \cdot \mathbf{v}_{is} \right) (-\mathbf{v}_{iq}) + \lambda \mathbf{u}_{uq} \quad \forall u \in \{1, \dots, m\}, q \in \{1, \dots, k\}$$

$$= \sum_{i:(u,i) \in \mathcal{S}} (e_{ui}) (-\mathbf{v}_{iq}) + \lambda \mathbf{u}_{uq} \quad \forall u \in \{1, \dots, m\}, q \in \{1, \dots, k\}$$

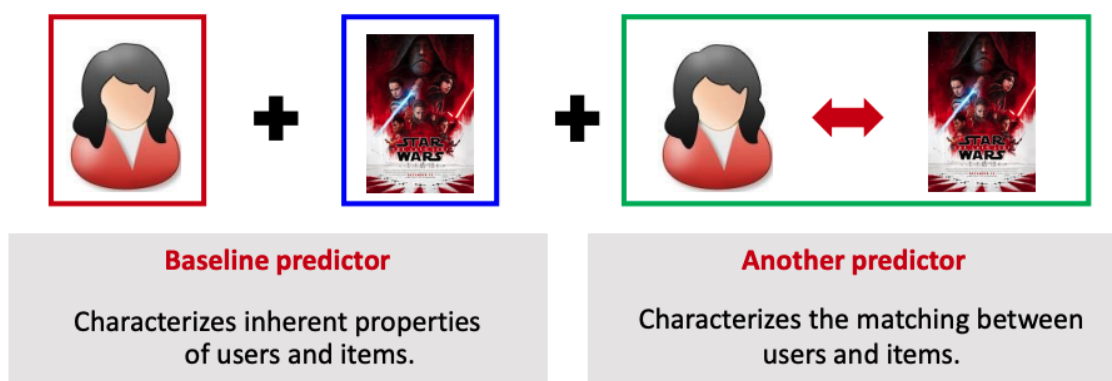
$$\frac{\partial E}{\partial \mathbf{v}_{iq}} = \sum_{u:(u,i) \in \mathcal{S}} \left(r_{ui} - \sum_{s=1}^k \mathbf{u}_{us} \cdot \mathbf{v}_{is} \right) (-\mathbf{u}_{uq}) + \lambda \mathbf{v}_{iq} \quad \forall i \in \{1, \dots, n\}, q \in \{1, \dots, k\}$$

$$= \sum_{u:(u,i) \in \mathcal{S}} (e_{ui}) (-\mathbf{u}_{uq}) + \lambda \mathbf{v}_{iq} \quad \forall i \in \{1, \dots, n\}, q \in \{1, \dots, k\}$$

2. Incorporating User and Item Biases

Modeling User Bias and Item Bias(Biased MF)

- The user rating consists of four parts:
 - μ : **global average** for all ratings
 - o_u : **user bias** for user ratings, p_i : **item bias** for item ratings
 - $u_u v_i^T$: **interaction** between user u and item i



➤ **Mean rating:** $\mu = 3.5$

➤ **Alice is a critical reviewer.**

- ◆ Your ratings are 1 star lower than the mean: $o_u = -1.0$

➤ **Star Wars is a popular movie.**

- ◆ This gets a rating of 0.5 higher than the average movie: $p_i = +0.5$

➤ **Predicted rating for Alice on Star Wars:** $3.5 - 1.0 + 0.5 = 3.0$

$$r_{ui} = \underbrace{\mu}_{\text{Overall mean rating}} + \underbrace{o_u}_{\text{Bias for Alice}} + \underbrace{p_i}_{\text{Bias for Star Wars}} + \underbrace{u_u v_i^T}_{\text{Residual User-Movie interaction}}$$

Objective Function with Biases

- Assume that the matrix R is mean-centered by subtracting the global mean μ from the rating matrix.

Goodness of fit

$$\min \frac{1}{2} \sum_{(u,i) \in \mathcal{S}} \left(r_{ui} - (o_u + p_i + \mathbf{u}_u \mathbf{v}_i^T) \right)^2$$

Regularization

$$+ \frac{\lambda}{2} \left(\sum_{u=1}^m \|\mathbf{u}_u\|^2 + \sum_{i=1}^n \|\mathbf{v}_i\|^2 + \sum_{u=1}^m \|o_u\|^2 + \sum_{i=1}^n \|p_i\|^2 \right)$$


- Train with the SGD to find optimal parameters.

- ◆ o_u , p_i , \mathbf{u}_u and \mathbf{v}_i are treated as the parameters to be learned.

o_u : # user (m 个), p_i : # item (n 个), $\mathbf{u}_u \mathbf{v}_i^T$: $(m+n) \times k$

Tricks to Incorporate Biases

- Instead of having separate biases, we create larger factor matrices of size $m \times (k+2)$ and $n \times (k+2)$.
- ◆ Set the last column of the user factor matrix to all 1s.
- ◆ Set the second last column of the item factor matrix to all 1s.



Discussion



➤ Adding user and item biases can help reduce overfitting.

- ◆ It can improve the **generalizability of the learning algorithm** to unseen entries, especially for cold-start users and items.
- ◆ The **(non-personalized) predictions of using the item bias** can give **reasonable predictions**.



➤ Practical lessons learned from the Netflix Prize contest

“Of the numerous new algorithmic contributions, I would like to highlight one – those humble baseline predictors (or biases), which capture main effects in the data. While the literature mostly concentrates on the more sophisticated algorithmic aspects, we have learned that an accurate treatment of main effects is probably at least as significant as coming up with modeling breakthroughs.”

3. Incorporating Implicit Feedback

How to Derive Implicit Feedback?

- For explicit feedback, we can derive implicit feedback.
 - Implicit feedback is captured by the identity of the rated items, regardless of actual rating values.
- **What if using two different item factor matrices?**
 - We have **explicit** and **implicit** item factors.
 - **User factors** are derived as a **linear combination of implicit item factors**.

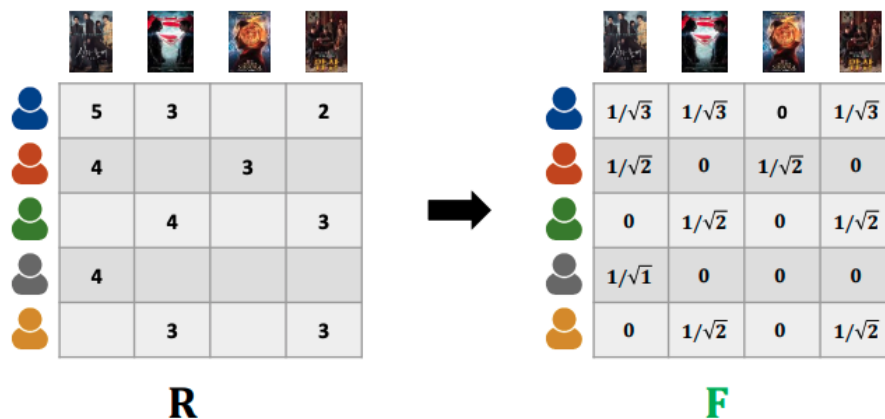
Buildin an Implicit Feedback Matrix

➤ An implicit feedback matrix is defined by the binary matrix.

- ◆ If a rating is observed, it is 1. otherwise, it is 0.


➤ Then, it is normalized so that the L2-norm of each row is 1.

- ◆ Each nonzero entry in the matrix \mathbf{F} is $1/\sqrt{|\mathcal{I}_u|}$, where \mathcal{I}_u is a set of items rated by user u .



Representing Explicit Feedback

- Assume that \mathbf{U} is computed by a combination of implicit matrix \mathbf{F} and implicit latent item matrix \mathbf{Y} .
 - The variables in \mathbf{Y} encode the propensity of each factor item to contribute to implicit feedback.



Asymmetric Factor Models

➤ How do we compute \hat{r}_{ui} ?

$$\hat{r}_{ui} = \sum_{s=1}^k [\mathbf{FY}]_{us} \cdot v_{is}$$

➤ Objective function

$$\min \frac{1}{2} \sum_{(u,i) \in \mathcal{S}} (r_{ui} - \hat{r}_{ui})^2 = \frac{1}{2} \sum_{(u,i) \in \mathcal{S}} \left(r_{ui} - \sum_{s=1}^k [\mathbf{FY}]_{us} \cdot v_{is} \right)^2$$

- ◆ It has two parameters \mathbf{Y} and \mathbf{V} .
- ◆ It is trained with the gradient descent method.

Pros : Asymmetric Factor Models

- It often **provides better results** than existing latent factor models.
 - Two users will have **similar user factors** if **they have rated similar items**, regardless of their rating values.
 - It can **reduce the redundancy in user factors** by deriving them as linear combinations of item factors.
- It also supports explainability.
 - We rewrite the factorization $[\mathbf{FY}]V^T$ as $F[\mathbf{Y}V^T]$.
 - The item-to-item prediction matrix $[\mathbf{Y}V^T]_{ij}$ tell us how much the act of rating item i contributes to the predicted rating of item j .
 - This type of explainability is inherent to **item-centric models**.

Cons: Asymmetric Factor Models



➤ Deriving user factors from the identities of rated items **may be an extreme case of using implicit feedback.**

- ◆ It does not discriminate between two users who have rated the same set of items but have very different ratings.

➤ Solution: the implicit user factor matrix **FY** is **only used to adjust the explicit user factor matrix U.**

- ◆ That is, **FY** is added to **U** before multiplying with **V^T**.

$$\mathbf{R} \approx (\mathbf{U} + \mathbf{FY})\mathbf{V}^T$$

Explicit user factors Implicit user factors



SVD++(=Biased MF + Asymmetric Factor Model)

- It can be view as combining the **unconstrained matrix factorization model** and the **asymmetric factorization model**.
 - The term **SVD++** is **slightly misleading** because the basis vectors **are not orthogonal**. It lossely implies latent factor models.

➤ How do we compute \hat{r}_{ui} ?

- ◆ Note: it contains user and item biases in **U** and **V**.

$$\hat{r}_{ui} = \sum_{s=1}^{k+2} (u_{us} + [\mathbf{FY}]_{us}) \cdot v_{is} = \sum_{s=1}^{k+2} \left(u_{us} + \sum_{h \in \mathcal{I}_u} \frac{y_{hs}}{\sqrt{J_u}} \right) \cdot v_{is}$$

\mathcal{I}_u : a set of items rated by user u

- ◆ The first term is computed by **UV^T**.
- ◆ The second term is computed by **[FY]V^T**.

Objective Function for SVD++

➤ SVD++ has three parameters **U**, **V**, and **Y** with biases.

Goodness of fit

$$\min \frac{1}{2} \sum_{(u,i) \in \mathcal{S}} \left(r_{ui} - \sum_{s=1}^{k+2} \left(\mathbf{u}_{us} + \sum_{h \in \mathcal{I}_u} \frac{y_{hs}}{\sqrt{J_u}} \right) \cdot \mathbf{v}_{is} \right)^2$$

Regularization

$$+ \frac{\lambda}{2} \sum_{s=1}^{k+1} \left(\sum_{u=1}^m \mathbf{u}_{us}^2 + \sum_{i=1}^n \mathbf{v}_{is}^2 + \sum_{i=1}^n \mathbf{y}_{is}^2 \right)$$

- ◆ $(k + 2)$ th column of **U** contains only 1s.
- ◆ $(k + 1)$ th column of **V** contains only 1s.
- ◆ Last two columns of **Y** contain only 0s.

Training with SGD



➤ Use the partial derivative to derive update rules.

$$\begin{aligned} \hat{r}_{ui} &= \sum_{s=1}^{k+2} \left(\mathbf{u}_{us} + \sum_{h \in \mathcal{I}_u} \frac{y_{hs}}{\sqrt{J_u}} \right) \cdot \mathbf{v}_{is} \\ e_{ui} &= r_{ui} - \hat{r}_{ui} \\ \mathbf{u}_{uq} &\leftarrow \mathbf{u}_{uq} + \alpha (e_{ui} \cdot \mathbf{v}_{iq} - \lambda \cdot \mathbf{u}_{uq}) \quad \forall q \in \{1, \dots, k+2\} \\ \mathbf{v}_{iq} &\leftarrow \mathbf{v}_{iq} + \alpha \left(e_{ui} \cdot \left(\mathbf{u}_{us} + \sum_{h \in \mathcal{I}_u} \frac{y_{hs}}{\sqrt{J_u}} \right) - \lambda \cdot \mathbf{v}_{iq} \right) \quad \forall q \in \{1, \dots, k+2\} \\ y_{hq} &\leftarrow y_{hq} + \alpha \left(\frac{e_{ui} \cdot \mathbf{v}_{iq}}{\sqrt{J_u}} - \lambda \cdot y_{hq} \right) \quad \forall q \in \{1, \dots, k+2\}, \forall h \in \mathcal{I}_u \end{aligned}$$

➤ Update **U**, **V**, and **Y** with biases with SGD.

Training with SGD



➤ It can be rewritten in vectorized form.

$$\begin{aligned}\mathbf{u}_u &\leftarrow \mathbf{u}_u + \alpha(e_{ui}\mathbf{v}_i - \lambda\mathbf{u}_u) \\ \mathbf{v}_i &\leftarrow \mathbf{v}_i + \alpha\left(e_{ui} \cdot \left(u_{us} + \sum_{h \in \mathcal{I}_u} \frac{y_{hs}}{\sqrt{\mathcal{I}_u}}\right) - \lambda\mathbf{v}_i\right) \\ \mathbf{y}_h &\leftarrow \mathbf{y}_h + \alpha\left(\frac{e_{ui} \cdot \mathbf{v}_i}{\sqrt{\mathcal{I}_u}} - \lambda\mathbf{y}_h\right) \quad \forall h \in \mathcal{I}_u\end{aligned}$$

➤ These updates may be applied to the rows of **U**, **V**, and **Y**.

